

# Potato Disease Classification

## Import all the Dependencies

```
In [1]: import tensorflow as tf
        from tensorflow.keras import models, layers
        import matplotlib.pyplot as plt
        from IPython.display import HTML
```

WARNING:tensorflow:From C:\Users\91760\anaconda3\lib\site-packages\keras\src\losses.py:2976: The name tf.losses.sparse\_softmax\_cross\_entropy is deprecated. Please use tf.compat.v1.losses.sparse\_softmax\_cross\_entropy instead.

## Set all the Constants

```
In [2]: BATCH_SIZE = 32
        IMAGE_SIZE = 256
        CHANNELS=3
        EPOCHS=50
```

## Import data into tensorflow dataset object

```
In [3]: dataset = tf.keras.preprocessing.image_dataset_from_directory(
        "training",
        seed=123,
        shuffle=True,
        image_size=(IMAGE_SIZE, IMAGE_SIZE),
        batch_size=BATCH_SIZE
    )
```

Found 2152 files belonging to 3 classes.

```
In [4]: class_names = dataset.class_names
        class_names
```

```
Out[4]: ['Potato__Early_blight', 'Potato__Late_blight', 'Potato__healthy']
```

```
In [5]: for image_batch, labels_batch in dataset.take(1):
        print(image_batch.shape)
        print(labels_batch.numpy())
```

```
(32, 256, 256, 3)
[1 1 1 0 0 0 0 0 1 1 1 0 1 0 1 1 1 0 1 0 1 0 0 1 0 0 1 1 2 0 0]
```

## Visualize some of the images from dataset

```
In [6]: plt.figure(figsize=(10, 10))
        for image_batch, labels_batch in dataset.take(1):
            for i in range(12):
                ax = plt.subplot(3, 4, i + 1)
                plt.imshow(image_batch[i].numpy().astype("uint8"))
                plt.title(class_names[labels_batch[i]])
                plt.axis("off")
```

Potato\_\_Early\_blight



Potato\_\_Early\_blight



Potato\_\_Early\_blight



Potato\_\_Late\_blight



Potato\_\_Early\_blight



Potato\_\_Early\_blight



Potato\_\_Late\_blight



Potato\_\_Early\_blight



Potato\_\_Late\_blight



Potato\_\_Early\_blight



Potato\_\_Early\_blight



Potato\_\_Early\_blight



## Function to Split Dataset

```
In [7]: len(dataset)
```

```
Out[7]: 68
```

```
In [8]: train_size = 0.8
len(dataset)*train_size
```

```
Out[8]: 54.400000000000006
```

```
In [9]: train_ds = dataset.take(54)
len(train_ds)
```

```
Out[9]: 54
```

```
In [10]: test_ds = dataset.skip(54)
len(test_ds)
```

```
Out[10]: 14
```

```
In [11]: val_size=0.1
len(dataset)*val_size
```

```
Out[11]: 6.800000000000001
```

```
In [12]: val_ds = test_ds.take(6)
len(val_ds)
```

```
Out[12]: 6
```

```
In [13]: test_ds = test_ds.skip(6)
len(test_ds)
```

```
Out[13]: 8
```

```
In [14]: def get_dataset_partitions_tf(ds, train_split=0.8, val_split=0.1, test_split=0.1, shuffle=True):
    assert (train_split + test_split + val_split) == 1

    ds_size = len(ds)

    if shuffle:
        ds = ds.shuffle(shuffle_size, seed=12)

    train_size = int(train_split * ds_size)
    val_size = int(val_split * ds_size)

    train_ds = ds.take(train_size)
    val_ds = ds.skip(train_size).take(val_size)
    test_ds = ds.skip(train_size).skip(val_size)

    return train_ds, val_ds, test_ds
```

```
In [15]: train_ds, val_ds, test_ds = get_dataset_partitions_tf(dataset)
```

```
In [16]: len(train_ds)
```

```
Out[16]: 54
```

```
In [17]: len(val_ds)
```

```
Out[17]: 6
```

```
In [18]: len(test_ds)
```

```
Out[18]: 8
```

## Cache, Shuffle, and Prefetch the Dataset

```
In [19]: train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=tf.data.AUTOTUNE)
val_ds = val_ds.cache().shuffle(1000).prefetch(buffer_size=tf.data.AUTOTUNE)
test_ds = test_ds.cache().shuffle(1000).prefetch(buffer_size=tf.data.AUTOTUNE)
```

## Building the Model

```
In [20]: resize_and_rescale = tf.keras.Sequential([
    layers.experimental.preprocessing.Resizing(IMAGE_SIZE, IMAGE_SIZE),
    layers.experimental.preprocessing.Rescaling(1./255),
])
```

WARNING:tensorflow:From C:\Users\91760\anaconda3\lib\site-packages\keras\src\backend.py:873: The name tf.get\_default\_graph is deprecated. Please use tf.compat.v1.get\_default\_graph instead.

## Data Augmentation

```
In [21]: data_augmentation = tf.keras.Sequential([
        layers.experimental.preprocessing.RandomFlip("horizontal_and_vertical"),
        layers.experimental.preprocessing.RandomRotation(0.2),
    ])
```

## Applying Data Augmentation to Train Dataset

```
In [22]: train_ds = train_ds.map(
        lambda x, y: (data_augmentation(x, training=True), y)
    ).prefetch(buffer_size=tf.data.AUTOTUNE)
```

## Model Architecture

```
In [23]: input_shape = (BATCH_SIZE, IMAGE_SIZE, IMAGE_SIZE, CHANNELS)
        n_classes = 3

        model = models.Sequential([
            resize_and_rescale,
            layers.Conv2D(32, kernel_size = (3,3), activation='relu', input_shape=input_shape),
            layers.MaxPooling2D((2, 2)),
            layers.Conv2D(64, kernel_size = (3,3), activation='relu'),
            layers.MaxPooling2D((2, 2)),
            layers.Conv2D(64, kernel_size = (3,3), activation='relu'),
            layers.MaxPooling2D((2, 2)),
            layers.Conv2D(64, (3, 3), activation='relu'),
            layers.MaxPooling2D((2, 2)),
            layers.Conv2D(64, (3, 3), activation='relu'),
            layers.MaxPooling2D((2, 2)),
            layers.Conv2D(64, (3, 3), activation='relu'),
            layers.MaxPooling2D((2, 2)),
            layers.Flatten(),
            layers.Dense(64, activation='relu'),
            layers.Dense(n_classes, activation='softmax'),
        ])

        model.build(input_shape=input_shape)
```

WARNING:tensorflow:From C:\Users\91760\anaconda3\lib\site-packages\keras\src\layers\pooling\max\_pooling2d.py:161: The name tf.nn.max\_pool is deprecated. Please use tf.nn.max\_pool2d instead.

```
In [24]: model.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
=====		
sequential (Sequential)	(32, 256, 256, 3)	0
conv2d (Conv2D)	(32, 254, 254, 32)	896
max_pooling2d (MaxPooling2D)	(32, 127, 127, 32)	0
conv2d_1 (Conv2D)	(32, 125, 125, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(32, 62, 62, 64)	0
conv2d_2 (Conv2D)	(32, 60, 60, 64)	36928
max_pooling2d_2 (MaxPooling2D)	(32, 30, 30, 64)	0
conv2d_3 (Conv2D)	(32, 28, 28, 64)	36928

max_pooling2d_3 (MaxPoolin g2D)	(32, 14, 14, 64)	0
conv2d_4 (Conv2D)	(32, 12, 12, 64)	36928
max_pooling2d_4 (MaxPoolin g2D)	(32, 6, 6, 64)	0
conv2d_5 (Conv2D)	(32, 4, 4, 64)	36928
max_pooling2d_5 (MaxPoolin g2D)	(32, 2, 2, 64)	0
flatten (Flatten)	(32, 256)	0
dense (Dense)	(32, 64)	16448
dense_1 (Dense)	(32, 3)	195

```

=====
Total params: 183747 (717.76 KB)
Trainable params: 183747 (717.76 KB)
Non-trainable params: 0 (0.00 Byte)

```

---

## Compiling the Model

```
In [25]: model.compile(
          optimizer='adam',
          loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False),
          metrics=['accuracy']
        )
```

```

WARNING:tensorflow:From C:\Users\91760\anaconda3\lib\site-packages\keras\src\optimize
rs\__init__.py:309: The name tf.train.Optimizer is deprecated. Please use tf.compat.v
1.train.Optimizer instead.

```

```
In [26]: history = model.fit(
          train_ds,
          batch_size=BATCH_SIZE,
          validation_data=val_ds,
          verbose=1,
          epochs=50,
        )
```

Epoch 1/50

```

WARNING:tensorflow:From C:\Users\91760\anaconda3\lib\site-packages\keras\src\utils\tf
_utils.py:492: The name tf.ragged.RaggedTensorValue is deprecated. Please use tf.comp
at.v1.ragged.RaggedTensorValue instead.

```

```

WARNING:tensorflow:From C:\Users\91760\anaconda3\lib\site-packages\keras\src\engine\b
ase_layer_utils.py:384: The name tf.executing_eagerly_outside_functions is deprecate
d. Please use tf.compat.v1.executing_eagerly_outside_functions instead.

```

```

54/54 [=====] - 251s 4s/step - loss: 0.8926 - accuracy: 0.51
62 - val_loss: 0.7564 - val_accuracy: 0.6406

```

Epoch 2/50

```

54/54 [=====] - 220s 4s/step - loss: 0.5741 - accuracy: 0.72
63 - val_loss: 0.4077 - val_accuracy: 0.8125

```

Epoch 3/50

```

54/54 [=====] - 1391s 26s/step - loss: 0.3620 - accuracy: 0.
8513 - val_loss: 0.3881 - val_accuracy: 0.8229

```

Epoch 4/50

```

54/54 [=====] - 216s 4s/step - loss: 0.2209 - accuracy: 0.90
45 - val_loss: 0.2484 - val_accuracy: 0.8854

```

Epoch 5/50

54/54 [=====] - 213s 4s/step - loss: 0.1090 - accuracy: 0.94  
91 - val\_loss: 0.1930 - val\_accuracy: 0.9323  
Epoch 6/50  
54/54 [=====] - 1020s 19s/step - loss: 0.1622 - accuracy: 0.  
9392 - val\_loss: 0.1460 - val\_accuracy: 0.9479  
Epoch 7/50  
54/54 [=====] - 213s 4s/step - loss: 0.1117 - accuracy: 0.96  
06 - val\_loss: 0.0626 - val\_accuracy: 0.9740  
Epoch 8/50  
54/54 [=====] - 210s 4s/step - loss: 0.0660 - accuracy: 0.97  
69 - val\_loss: 0.2849 - val\_accuracy: 0.8958  
Epoch 9/50  
54/54 [=====] - 212s 4s/step - loss: 0.0751 - accuracy: 0.97  
40 - val\_loss: 0.1044 - val\_accuracy: 0.9479  
Epoch 10/50  
54/54 [=====] - 211s 4s/step - loss: 0.0815 - accuracy: 0.97  
11 - val\_loss: 0.1703 - val\_accuracy: 0.9427  
Epoch 11/50  
54/54 [=====] - 212s 4s/step - loss: 0.0646 - accuracy: 0.97  
57 - val\_loss: 0.1237 - val\_accuracy: 0.9427  
Epoch 12/50  
54/54 [=====] - 221s 4s/step - loss: 0.0960 - accuracy: 0.96  
24 - val\_loss: 0.0269 - val\_accuracy: 0.9948  
Epoch 13/50  
54/54 [=====] - 214s 4s/step - loss: 0.0543 - accuracy: 0.97  
86 - val\_loss: 0.0261 - val\_accuracy: 0.9896  
Epoch 14/50  
54/54 [=====] - 212s 4s/step - loss: 0.0508 - accuracy: 0.97  
80 - val\_loss: 0.2158 - val\_accuracy: 0.9323  
Epoch 15/50  
54/54 [=====] - 211s 4s/step - loss: 0.0368 - accuracy: 0.98  
73 - val\_loss: 0.0639 - val\_accuracy: 0.9635  
Epoch 16/50  
54/54 [=====] - 171s 3s/step - loss: 0.0516 - accuracy: 0.98  
03 - val\_loss: 0.8642 - val\_accuracy: 0.8281  
Epoch 17/50  
54/54 [=====] - 176s 3s/step - loss: 0.0932 - accuracy: 0.96  
64 - val\_loss: 0.0656 - val\_accuracy: 0.9635  
Epoch 18/50  
54/54 [=====] - 178s 3s/step - loss: 0.0403 - accuracy: 0.98  
61 - val\_loss: 0.0632 - val\_accuracy: 0.9792  
Epoch 19/50  
54/54 [=====] - 190s 4s/step - loss: 0.0280 - accuracy: 0.99  
02 - val\_loss: 0.0261 - val\_accuracy: 0.9792  
Epoch 20/50  
54/54 [=====] - 184s 3s/step - loss: 0.0372 - accuracy: 0.98  
90 - val\_loss: 0.0876 - val\_accuracy: 0.9635  
Epoch 21/50  
54/54 [=====] - 173s 3s/step - loss: 0.0514 - accuracy: 0.98  
03 - val\_loss: 0.1123 - val\_accuracy: 0.9635  
Epoch 22/50  
54/54 [=====] - 177s 3s/step - loss: 0.0450 - accuracy: 0.98  
15 - val\_loss: 0.0742 - val\_accuracy: 0.9531  
Epoch 23/50  
54/54 [=====] - 177s 3s/step - loss: 0.0896 - accuracy: 0.96  
93 - val\_loss: 0.0531 - val\_accuracy: 0.9688  
Epoch 24/50  
54/54 [=====] - 159s 3s/step - loss: 0.0237 - accuracy: 0.99  
25 - val\_loss: 0.0054 - val\_accuracy: 1.0000  
Epoch 25/50  
54/54 [=====] - 148s 3s/step - loss: 0.0238 - accuracy: 0.99  
31 - val\_loss: 0.0289 - val\_accuracy: 0.9844  
Epoch 26/50  
54/54 [=====] - 152s 3s/step - loss: 0.0530 - accuracy: 0.98  
50 - val\_loss: 0.5365 - val\_accuracy: 0.8646  
Epoch 27/50  
54/54 [=====] - 154s 3s/step - loss: 0.0415 - accuracy: 0.98  
55 - val\_loss: 0.0485 - val\_accuracy: 0.9740  
Epoch 28/50

54/54 [=====] - 148s 3s/step - loss: 0.0325 - accuracy: 0.98  
78 - val\_loss: 0.1546 - val\_accuracy: 0.9479  
Epoch 29/50  
54/54 [=====] - 193s 4s/step - loss: 0.0397 - accuracy: 0.98  
55 - val\_loss: 0.1268 - val\_accuracy: 0.9635  
Epoch 30/50  
54/54 [=====] - 198s 4s/step - loss: 0.0288 - accuracy: 0.98  
96 - val\_loss: 0.0237 - val\_accuracy: 0.9948  
Epoch 31/50  
54/54 [=====] - 165s 3s/step - loss: 0.0205 - accuracy: 0.99  
36 - val\_loss: 0.0031 - val\_accuracy: 1.0000  
Epoch 32/50  
54/54 [=====] - 193s 4s/step - loss: 0.0363 - accuracy: 0.98  
44 - val\_loss: 0.1079 - val\_accuracy: 0.9583  
Epoch 33/50  
54/54 [=====] - 157s 3s/step - loss: 0.0338 - accuracy: 0.98  
84 - val\_loss: 0.1397 - val\_accuracy: 0.9583  
Epoch 34/50  
54/54 [=====] - 155s 3s/step - loss: 0.0396 - accuracy: 0.98  
67 - val\_loss: 0.0273 - val\_accuracy: 0.9844  
Epoch 35/50  
54/54 [=====] - 150s 3s/step - loss: 0.0358 - accuracy: 0.98  
96 - val\_loss: 0.0297 - val\_accuracy: 0.9844  
Epoch 36/50  
54/54 [=====] - 151s 3s/step - loss: 0.0533 - accuracy: 0.98  
09 - val\_loss: 0.0136 - val\_accuracy: 0.9948  
Epoch 37/50  
54/54 [=====] - 150s 3s/step - loss: 0.0377 - accuracy: 0.98  
73 - val\_loss: 0.0376 - val\_accuracy: 0.9792  
Epoch 38/50  
54/54 [=====] - 150s 3s/step - loss: 0.0176 - accuracy: 0.99  
42 - val\_loss: 0.0511 - val\_accuracy: 0.9740  
Epoch 39/50  
54/54 [=====] - 149s 3s/step - loss: 0.0152 - accuracy: 0.99  
48 - val\_loss: 0.0268 - val\_accuracy: 0.9896  
Epoch 40/50  
54/54 [=====] - 147s 3s/step - loss: 0.0264 - accuracy: 0.99  
07 - val\_loss: 0.0479 - val\_accuracy: 0.9740  
Epoch 41/50  
54/54 [=====] - 147s 3s/step - loss: 0.0145 - accuracy: 0.99  
77 - val\_loss: 0.0143 - val\_accuracy: 0.9948  
Epoch 42/50  
54/54 [=====] - 147s 3s/step - loss: 0.0158 - accuracy: 0.99  
48 - val\_loss: 0.0991 - val\_accuracy: 0.9583  
Epoch 43/50  
54/54 [=====] - 151s 3s/step - loss: 0.0567 - accuracy: 0.98  
32 - val\_loss: 0.0796 - val\_accuracy: 0.9635  
Epoch 44/50  
54/54 [=====] - 149s 3s/step - loss: 0.0299 - accuracy: 0.98  
84 - val\_loss: 0.0452 - val\_accuracy: 0.9688  
Epoch 45/50  
54/54 [=====] - 151s 3s/step - loss: 0.0247 - accuracy: 0.99  
07 - val\_loss: 0.0528 - val\_accuracy: 0.9740  
Epoch 46/50  
54/54 [=====] - 166s 3s/step - loss: 0.0421 - accuracy: 0.98  
61 - val\_loss: 0.0463 - val\_accuracy: 0.9896  
Epoch 47/50  
54/54 [=====] - 154s 3s/step - loss: 0.0231 - accuracy: 0.99  
25 - val\_loss: 0.0148 - val\_accuracy: 0.9948  
Epoch 48/50  
54/54 [=====] - 160s 3s/step - loss: 0.0463 - accuracy: 0.98  
38 - val\_loss: 0.2533 - val\_accuracy: 0.9219  
Epoch 49/50  
54/54 [=====] - 152s 3s/step - loss: 0.0301 - accuracy: 0.98  
78 - val\_loss: 0.2292 - val\_accuracy: 0.9219  
Epoch 50/50  
54/54 [=====] - 151s 3s/step - loss: 0.0183 - accuracy: 0.99  
42 - val\_loss: 0.0176 - val\_accuracy: 0.9948



```

In [27]: scores = model.evaluate(test_ds)

8/8 [=====] - 10s 622ms/step - loss: 0.0389 - accuracy: 0.9844

In [28]: scores

Out[28]: [0.03888298571109772, 0.984375]

In [29]: history

Out[29]: <keras.src.callbacks.History at 0x1c92f96f910>

In [30]: history.params

Out[30]: {'verbose': 1, 'epochs': 50, 'steps': 54}

In [31]: history.history.keys()

Out[31]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])

In [32]: type(history.history['loss'])

Out[32]: list

In [33]: len(history.history['loss'])

Out[33]: 50

In [34]: history.history['loss'][:5] # show loss for first 5 epochs

Out[34]: [0.8925877809524536,
0.5740604400634766,
0.3619817793369293,
0.22088845074176788,
0.10896806418895721]

In [35]: acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

loss = history.history['loss']
val_loss = history.history['val_loss']


In [36]: plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(range(EPOCHS), acc, label='Training Accuracy')
plt.plot(range(EPOCHS), val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(range(EPOCHS), loss, label='Training Loss')
plt.plot(range(EPOCHS), val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()


```



Training and Validation Accuracy



Training and Validation Loss



## Run prediction on a sample image


```
In [37]: import numpy as np
for images_batch, labels_batch in test_ds.take(1):

    first_image = images_batch[0].numpy().astype('uint8')
    first_label = labels_batch[0].numpy()

    print("first image to predict")
    plt.imshow(first_image)
    print("actual label:", class_names[first_label])

    batch_prediction = model.predict(images_batch)
    print("predicted label:", class_names[np.argmax(batch_prediction[0])])
```

```
first image to predict
actual label: Potato__Early_blight
1/1 [=====] - 2s 2s/step
predicted label: Potato__Early_blight
```



Write a function for inference

```
In [38]: def predict(model, img):
    img_array = tf.keras.preprocessing.image.img_to_array(images[i].numpy())
    img_array = tf.expand_dims(img_array, 0)

    predictions = model.predict(img_array)

    predicted_class = class_names[np.argmax(predictions[0])]
    confidence = round(100 * (np.max(predictions[0])), 2)
    return predicted_class, confidence
```

Now run inference on few sample images

```
In [39]: plt.figure(figsize=(15, 15))
    for images, labels in test_ds.take(1):
        for i in range(9):
            ax = plt.subplot(3, 3, i + 1)
            plt.imshow(images[i].numpy().astype("uint8"))

            predicted_class, confidence = predict(model, images[i].numpy())
            actual_class = class_names[labels[i]]

            plt.title(f"Actual: {actual_class},\n Predicted: {predicted_class}.\n Confidence: {confidence}")

            plt.axis("off")
```

```
1/1 [=====] - 0s 301ms/step
1/1 [=====] - 0s 63ms/step
1/1 [=====] - 0s 63ms/step
1/1 [=====] - 0s 97ms/step
1/1 [=====] - 0s 69ms/step
1/1 [=====] - 0s 78ms/step
1/1 [=====] - 0s 49ms/step
1/1 [=====] - 0s 63ms/step
1/1 [=====] - 0s 50ms/step
```

Actual: Potato\_\_Late\_blight,  
Predicted: Potato\_\_Late\_blight.  
Confidence: 100.0%



Actual: Potato\_\_Late\_blight,  
Predicted: Potato\_\_Late\_blight.  
Confidence: 100.0%

Actual: Potato\_\_Early\_blight,  
Predicted: Potato\_\_Early\_blight.  
Confidence: 99.99%



Actual: Potato\_\_healthy,  
Predicted: Potato\_\_healthy.  
Confidence: 99.99%

Actual: Potato\_\_Late\_blight,  
Predicted: Potato\_\_Late\_blight.  
Confidence: 78.45%



Actual: Potato\_\_Late\_blight,  
Predicted: Potato\_\_Late\_blight.  
Confidence: 100.0%



Actual: Potato\_\_Late\_blight,  
Predicted: Potato\_\_Late\_blight.  
Confidence: 99.91%



Actual: Potato\_\_Late\_blight,  
Predicted: Potato\_\_Late\_blight.  
Confidence: 96.38%



Actual: Potato\_\_Early\_blight,  
Predicted: Potato\_\_Early\_blight.  
Confidence: 100.0%

